



Sistemas de Operação / Fundamentos de Sistemas Operativos I/O

Artur Pereira <artur@ua.pt>

DETI / Universidade de Aveiro

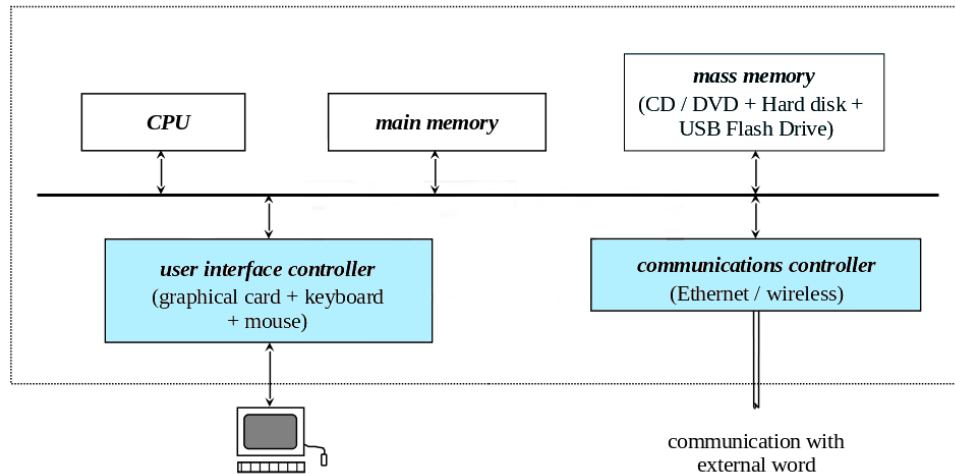
Outline

- 1 Overview
- 2 Device controllers
- 3 I/O programming
- 4 Types of access

Overview

The I/O components

- Simple view of a computational system, highlighting the I/O components




Overview

Role of the operating system

- Two distinct perspectives are usually considered for the role played by the operating system in managing the input / output devices of a computer system
 - **user perspective** – providing the application developer with a device interface (API) that is conceptually simple, reasonably uniform, and as much as possible independent of the specific device
 - **system perspective** – isolating the different devices from direct access by user processes by introducing a functional layer that directly controls the devices
 - send commands, transfer data, handle interrupts, and handle error conditions

Device controllers

Device and device controller



- There are 2 different components to consider
 - **The device itself** – physical system (electromechanical, optical-mechanical, ...), which stores information and converts it from, or to, an externally accessible form
 - **The device controller** – electronic circuit, more or less complex, part of the computational system, which works as an interface with the device
- From the point of view of the operating system, the **device controller** is the only relevant component
- Nowadays, controllers are very versatile, minimizing the role of the operating system in its management (programming)

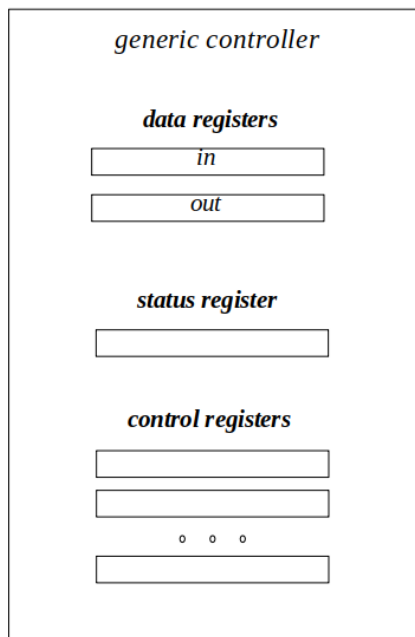
Device controllers

Types of devices

- In terms of transferring information, I/O devices fall into two broad categories
 - **character-type devices** – the transfer of information is based on a stream of bytes, whose length can be variable
 - **block-type devices** – the transfer of information is based on a constant and pre-defined number of bytes, the **block**, typically with a value equal to a power of 2 between 512 and 16K
- The way the transfer is done depends on the bus used
- The rate of transfer depends on the type of device
 - can vary from tens of bytes (keyboard, for example) to thousands of megabytes (SATA or USB3 disk)

Device controllers

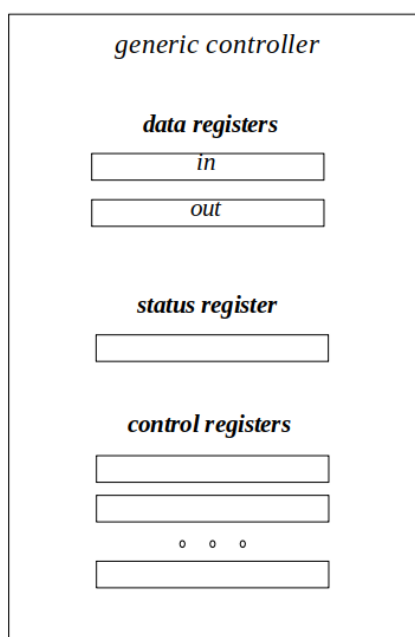
Internals



- A **generic controller**, from a programming point of view, can be seen as set of registers:
 - **control registers** – playing different functions
 - to configure the device
 - to define the type of interaction with the processor (polled I/O, interrupt-driven I/O or DMA-based I/O)
 - in complex controllers, to execute a command
 - a **status register** – representing the internal state of the device
 - to indicate the success of the last operation
 - to indicate the failure and errors of the last operation
 - to indicate it is ready to receive a new command

Device controllers


Internals (2)



- **data registers** – used for the communication itself
 - In **character-type** devices, the write and read commands are implicit
 - a value written in the **out** register is sent to the device
 - a value received from the device is put in the **in** register
 - In **block-type** devices, the transfer starts based on an explicit command
 - the data register is in general unique, **in-out**, and the direction of the transfer depends on the command given

Device controllers


Address modes



- There are 3 different possible ways to access the internal registers of a controller
 - **I/O-mapped** – controllers are mapped in a specific I/O address space
 - registers are accessed through specific instructions (**in** e **out**)
 - **memory-mapped I/O** – controllers are mapped in the memory address space
 - registers are accessed through the memory access instructions (**load** e **store**)
 - **hybrid** – controllers are mapped in a specific I/O address space, but data buffers are mapped in the memory address space to facilitate communication

Device controllers

Address modes (2)



- The Intel Pentium has an I/O address space of 64 KB
 - Some computer systems, based on Pentium, use this space to address controllers
 - But, a region in memory, between addresses 640 KB e 1 MB, is also reserved to implement data buffers for the devices

I/O programming

Objectives

- The environment provided by the operating system for the communication with the I/O devices should:
 - be independent of the device specifics
 - devices must be seen in a generic way
 - I/O redirecting, for example, should be possible in a natural way
 - support a uniform naming mechanism
 - device names must consist of strings of characters without any particular meaning
 - decouple devices from user processes
 - vast majority of I/O devices work in an asynchronous manner – data transfers to and from main memory are triggered by interrupts
 - from the user's perspective, however, it is simpler to design communication in a synchronous way – the process blocks until conditions are met for communication to take place


I/O programming

Objectives (2)

- The environment provided by the operating system for the communication with the I/O devices should (cont.):
 - manage access to preemptible and non-preemptible devices in a uniform way
 - communication with preemptible devices can be shared by multiple users simultaneously
 - communication with non-preemptible devices must take place in a mutual exclusion, or dedicated regime
 - the operating system therefore has to identify the different situations and ensure proper coordination
 - perform error management in an integrated manner
 - the detection of errors must be carried out as close to the device as possible in order to allow its [possible] recovery in a transparent way
 - the general policy should be to only report the error to the upper layer if the lower layer cannot handle it

Access types


Polled I/O



- In a **polled-I/O** approach there is no decoupling
 - It is the user process that is directly responsible for the communication
- **Device communication routines** are **system calls** that directly implement hardware access
- The **simplest solution**, but **little efficient**
 - The user processor enters a busy waiting, waiting for the completion of the operation

Access types

Polled I/O (2)



/ access routine; assumes a character-type device */*

```
void control(unsigned short add, unsigned char prog []);
```

```
#define RXRDY    ...    /* there are data to be read */
```

```
#define TXRDY    ...    /* transmitter register is empty */
```

```
#define ERROR    ...    /* error status */
```


```
unsigned char status(unsigned short add);
```

```
unsigned char in(unsigned short add);
```

```
void out(unsigned short add, unsigned char val);
```

Access types

Polled I/O (3)



/ possible routines for interaction with the device
system calls – running at kernel level
It is assumed that the communication channel was established already */*

*/**
\brief read N bytes
\param dd --- device descriptor, which represents the communication channel
\param N --- number of bytes to be read
\param buff --- pointer to storage area
\return 0 on success; -1 on error (errno is set accordingly)
/


```
int readNBytes (int dd, int N, unsigned char buff[]);
```

*/**
\brief write N bytes
\param dd --- device descriptor, which represents the communication channel
\param N --- number of bytes to be written
\param buff --- pointer to storage area
\return 0 on success; -1 on error (errno is set accordingly)
/

```
int writeNBytes (int dd, int N, unsigned char buff[]);
```

Access types

Polled I/O (4)




*/**
\param dd --- device descriptor, which represents the communication channel
\param N --- number of bytes to be read
\param buff --- pointer to storage area
\return 0 on success; -1 on error (errno is set accordingly)
/

```
int readNBytes(int dd, int N, unsigned char buff[])
```

```
{
    int add = getAdd(dd);
    for (int n = 0; n < N; n++)
    {
        int stat;
        do
        {
            stat = status(add);
            if ((stat & ERROR) != 0)
                /* error handling */
            } while ((stat & RXRDY) == 0);
            buff[n] = in(add);
        }
    }
    return 0;
}
```


Access types

Polled I/O (5)



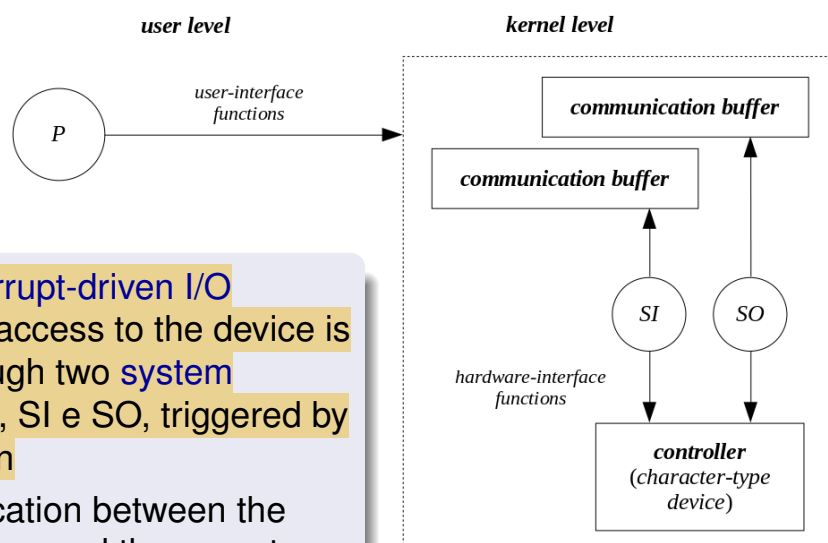
```

/**
 \param dd  --- device descriptor, which represents the communication channel
 \param N   --- number of bytes to be written
 \param buff --- pointer to storage area
 \return 0 on success; -1 on error (errno is set accordingly)
 */
int writeNBytes(int dd, int N, unsigned char buff [])
{
    int add = getAdd(dd);
    for (int n = 0; n < N; n++)
    {
        int stat;
        do
        {
            stat = status(add);
            if ((stat & ERROR) != 0)
                /* error handling */
            } while ((stat & TXRDY) == 0);
            out(add, buff[n]);
        }
    }
    return 0;
}

```

Access types

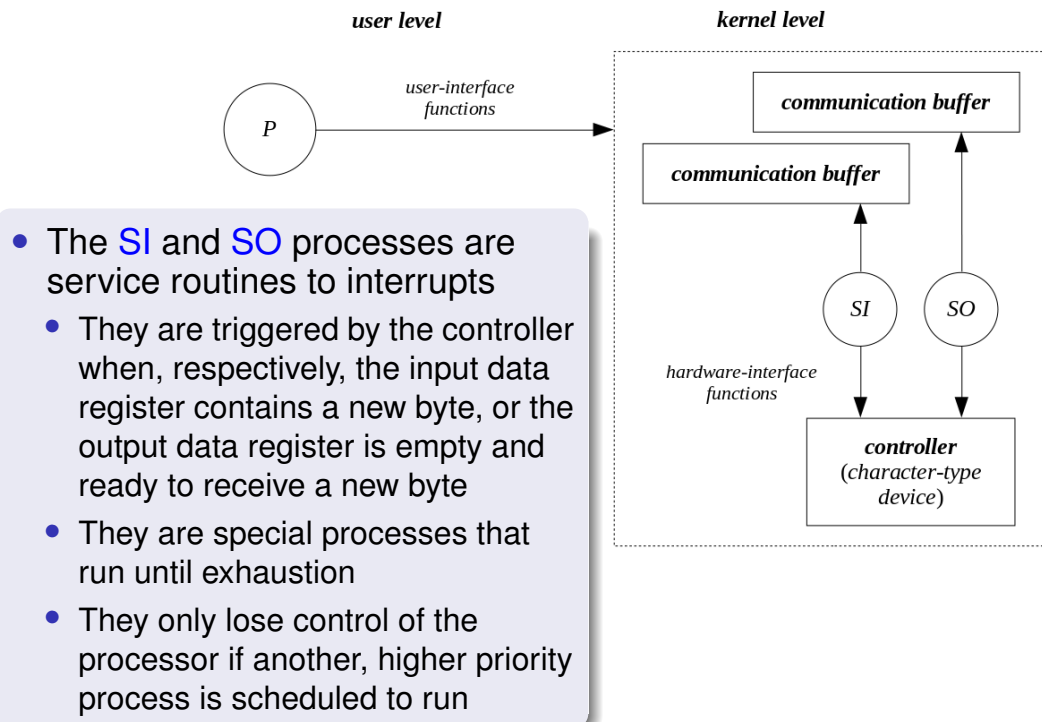
Interrupt-driven I/O



- In the **interrupt-driven I/O** approach access to the device is done through two **system processes**, **SI** e **SO**, triggered by interruption
- Communication between the user process and these system processes is done through **two communication buffers**

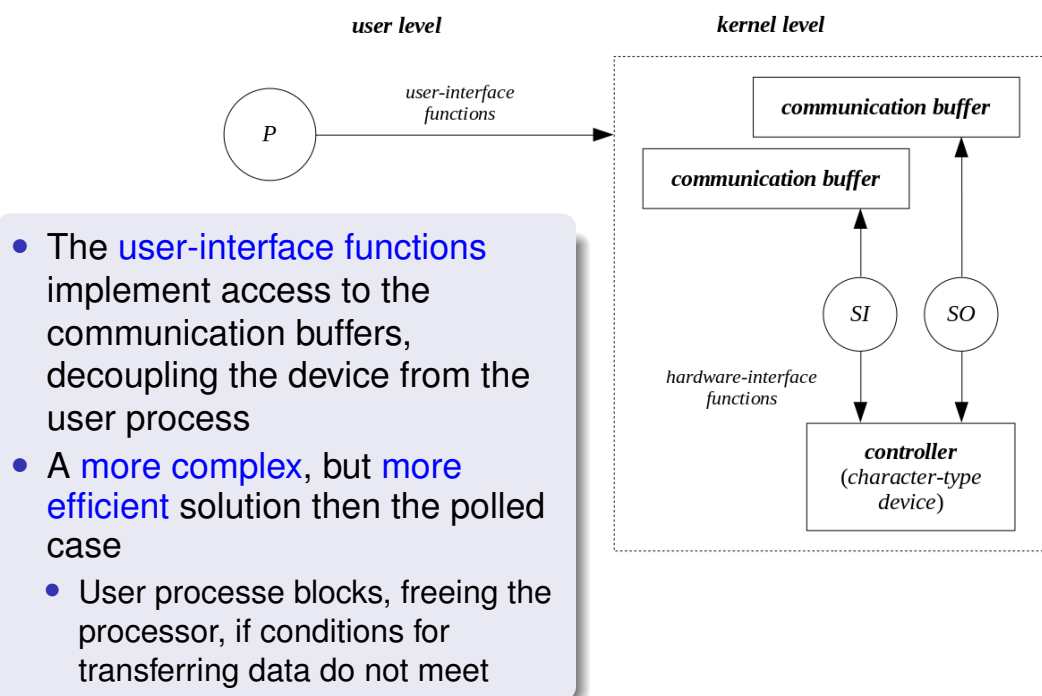
Access types

Interrupt-driven I/O (2)



Access types

Interrupt-driven I/O (3)



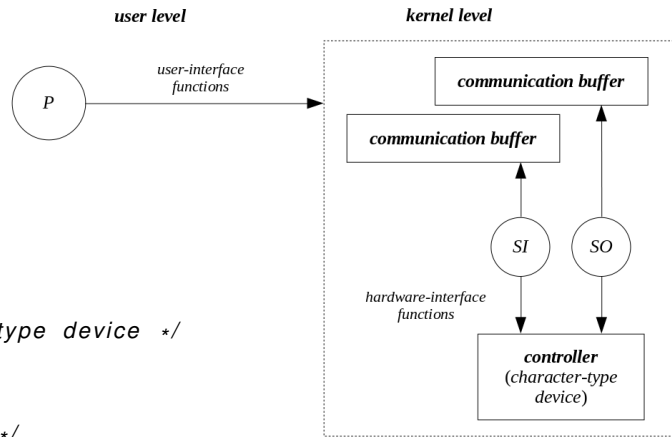
Access types

Interrupt-driven I/O (4)

/ access routine; assumes a character-type device */*

```
typedef struct COM.BUFF
{
    FIFO fifo;           /* storage area */
    SEMAPHORE wait;      /* user process' blocking semaphore */
    bool noMoreInt;      /* end of interruptions (only for output) */
    int errnum;          /* error status */
} COM.BUFF;
```

- in an **entry buffer**, the semaphore is initialized at 0
 - meaning that the FIFO is empty, so no data can be read
- in an **output buffer**, the semaphore is initialized at N, the size of the `fifo`
 - meaning that the FIFO is empty, so N bytes can be put there
- flag `noMoreInt` signals the need to prime the controller output data register so that interrupts are generated again
 - initialized to true



Access types

Interrupt-driven I/O (5)

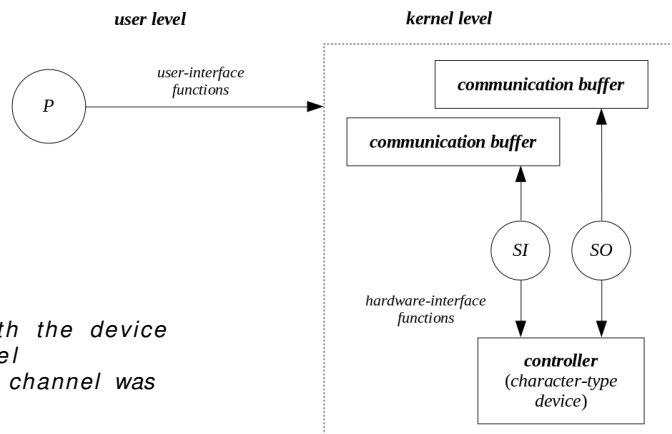
/ possible routines for interaction with the device
system calls – running at kernel level
It is assumed that the communication channel was
established already
/

```
/**
 * \brief read N bytes
 * \param dd --- device descriptor, which represents the communication channel
 * \param N --- number of bytes to be read
 * \param buff --- pointer to storage area
 * \return 0 on success; -1 on error (errno is set accordingly)
 */
```

```
int readNBytes (int dd, int N, unsigned char buff[]);
```

```
/**
 * \brief write N bytes
 * \param dd --- device descriptor, which represents the communication channel
 * \param N --- number of bytes to be written
 * \param buff --- pointer to storage area
 * \return 0 on success; -1 on error (errno is set accordingly)
 */
```

```
int writeNBytes (int dd, int N, unsigned char buff[]);
```



Access types

Interrupt-driven I/O (6)

```

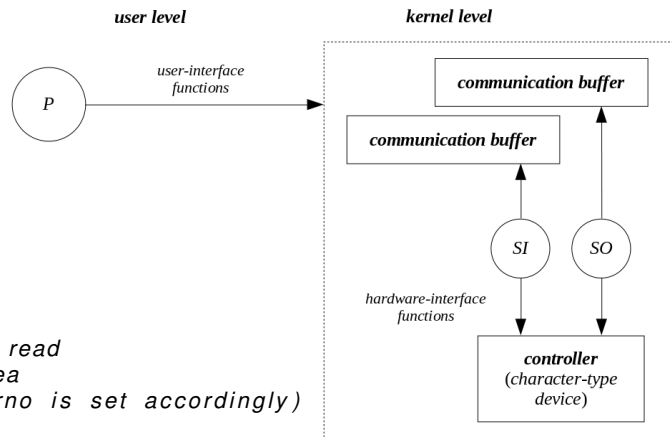
/**
 \brief read N bytes
 \param dd --- device descriptor
 \param N --- number of bytes to be read
 \param buff --- pointer to storage area
 \return 0 on success; -1 on error (errno is set accordingly)
 */
int readNBytes(int dd, int N, unsigned char buff[])
{
    COMM.BUF * inbuf = getBuff(dd);

    for (int n = 0; n < N; n++)
    {
        if (inbuf->errNumb != 0)
            /* error handling */

        sem_down(inbuf->wait);

        disable_interrupts();
        buff[n] = fifoOut(inbuf->fifo);
        enable_interrupts();
    }
    return 0;
}

```



Access types

Interrupt-driven I/O (7)

```

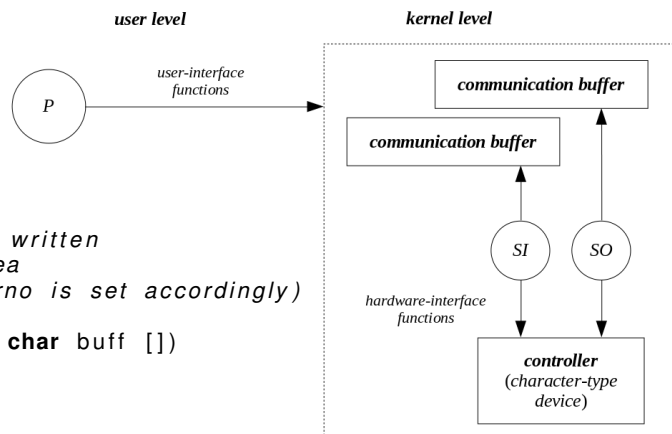
/**
 \brief write N bytes
 \param dd --- device descriptor
 \param N --- number of bytes to be written
 \param buff --- pointer to storage area
 \return 0 on success; -1 on error (errno is set accordingly)
 */
int writeNBytes(int dd, int N, unsigned char buff [])
{
    COMM.BUF * outbuf = getBuff(dd);

    for (int n = 0; n < N; n++)
    {
        if (outbuf->errNumb != 0)
            /* error handling */

        sem_down(outbuf->wait);

        disable_interrupts();
        fifoIn(outbuf->fifo, buff[n]);
        if (outbuf->noMoreInt)
        {
            writeDReg(getAdd(dd));
            Outbuf->noMoreInt = false;
        }
        enable_interrupts();
    }
    return 0;
}

```



Access types

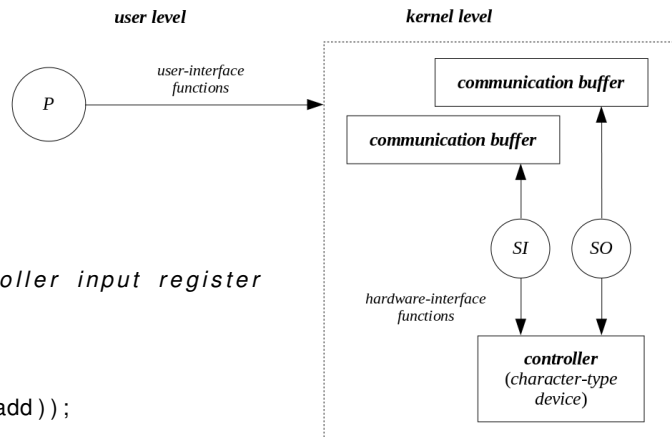
Interrupt-driven I/O (8)

```

/** SI process -
 * reading of one byte from the controller input register
 * add --- controller address
 */
void readDReg(unsigned short add)
{
    COM.BUFF * inbuf = getBuff(getDesc(add));

    int stat = status(add);          /* read the status register */
    if ((stat & ERROR) == ERROR)
        inBuf->errno = stat & ERROR; /* set error status */
    if ((stat & RXRDY) == RXRDY)     /* there are data to be read */
    {
        char val = in(add);
        if ((stat & ERROR) != ERROR)
        {
            if (not fifo_full(inbuf->fifo))
            {
                fifo_in(inbuf->fifo, val); /* store byte in buffer */
                sem_up(inbuf->wait);       /* signal there is data in buffer */
            }
            else
                inbuf->errNumb = OVERRUN; /* overrun error */
        }
    }
}

```



Access types

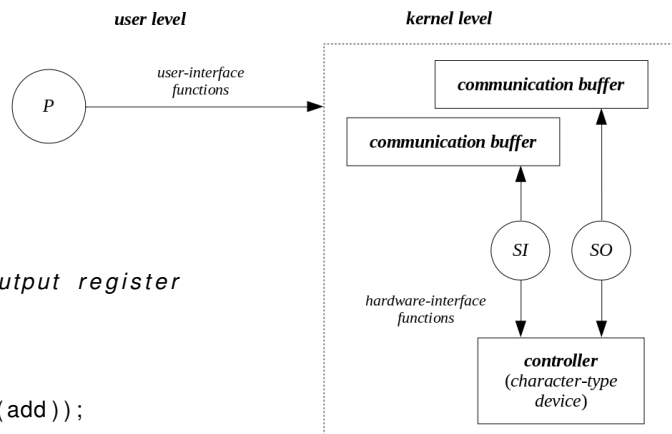
Interrupt-driven I/O (9)

```

/* SO process -
 * write one byte to the controller output register
 * add --- controller address
 */
void writeDReg(unsigned short add)
{
    COM.BUFF * outbuf = getBuff(getDesc(add));

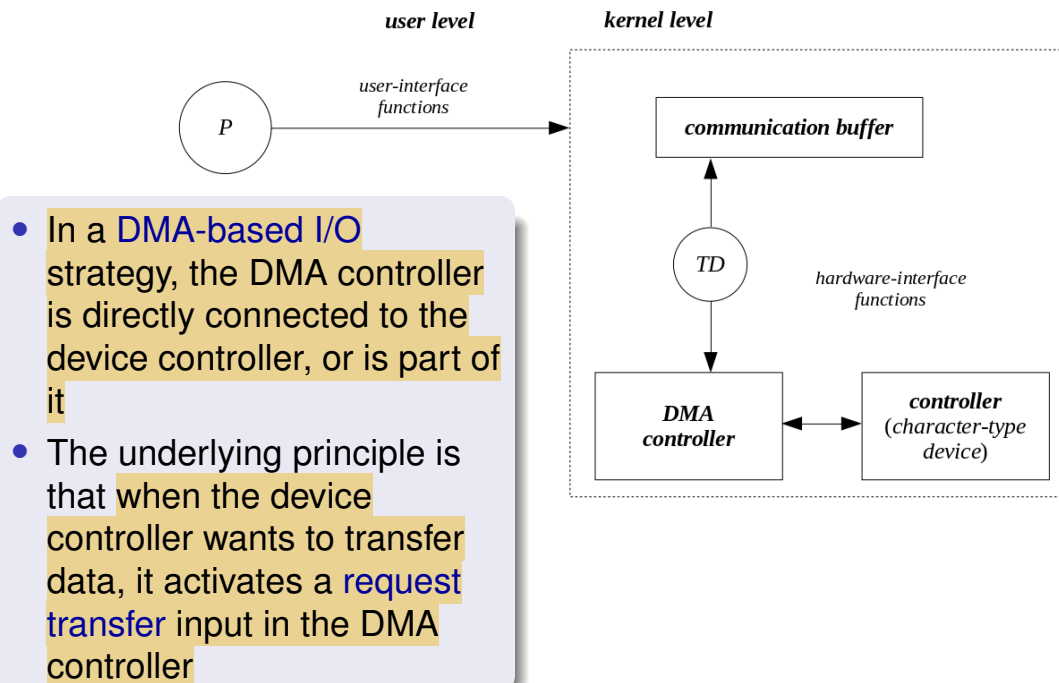
    int stat = status(add);          /* read status register */
    if ((stat & ERROR) == ERROR)
        outbuf->errno = stat & ERROR; /* set error status */
    if ((stat & TXRDY) == TXRDY)     /* can write ? */
    {
        if (not fifo_empty(outbuf->fifo))
        {
            char val = fifo_out(outbuf->fifo); /* retrieve byte from buffer */
            sem_up(outbuf->wait); /* signal there is room in buffer */
            out(add, val); /* put byte in controller output register */
        }
        else
            outbuf->noMoreInt = true; /* set end of interruptions */
    }
}

```



Access types

DMA-based I/O



Access types

DMA-based I/O (2)

